

# Introduction

This workshop is hosted in conjunction with OU's DALab and the Computer Science Graduate Student Association (CSGSA)

The notebook for this workshop and all workshops can be found on the [OU DALab github repo databitesp2020](#)

Presenter: Monique Shotande

## Would you like to be involved in Research at the University of Oklahoma?

This survey helps us understand the attendees' knowledge of Python and machine learning as well as expected outcomes from attending these sessions. We want to gauge demand of skills, resources, and programming knowledge. This information will allow us to continue to improve these workshops to meet your needs.

[Pre-knowledge Survey](#)

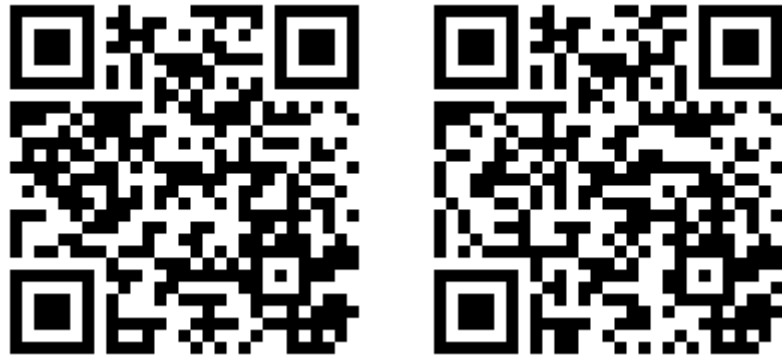


Please complete this brief [sign in form](#)

Check us out on social media

Facebook [@oucsgsa](#)

Instagram



## Autoencoders

- **Unsupervised** Neural Network Models
- Automate construction of optimal **compressed data representations**

## Categories of Machine Learning

- Supervised Learning
  - Features\* are matched with coresponding labels
- **Unsupervised Learning**
  - Only features are available
- Reinforcement Learning
  - Initial rules of engagement and reward system are established
  - The model updates rules based on how to achieve reward

\* piece of information describing the data

## Unsupervised Learning

Only the features are available

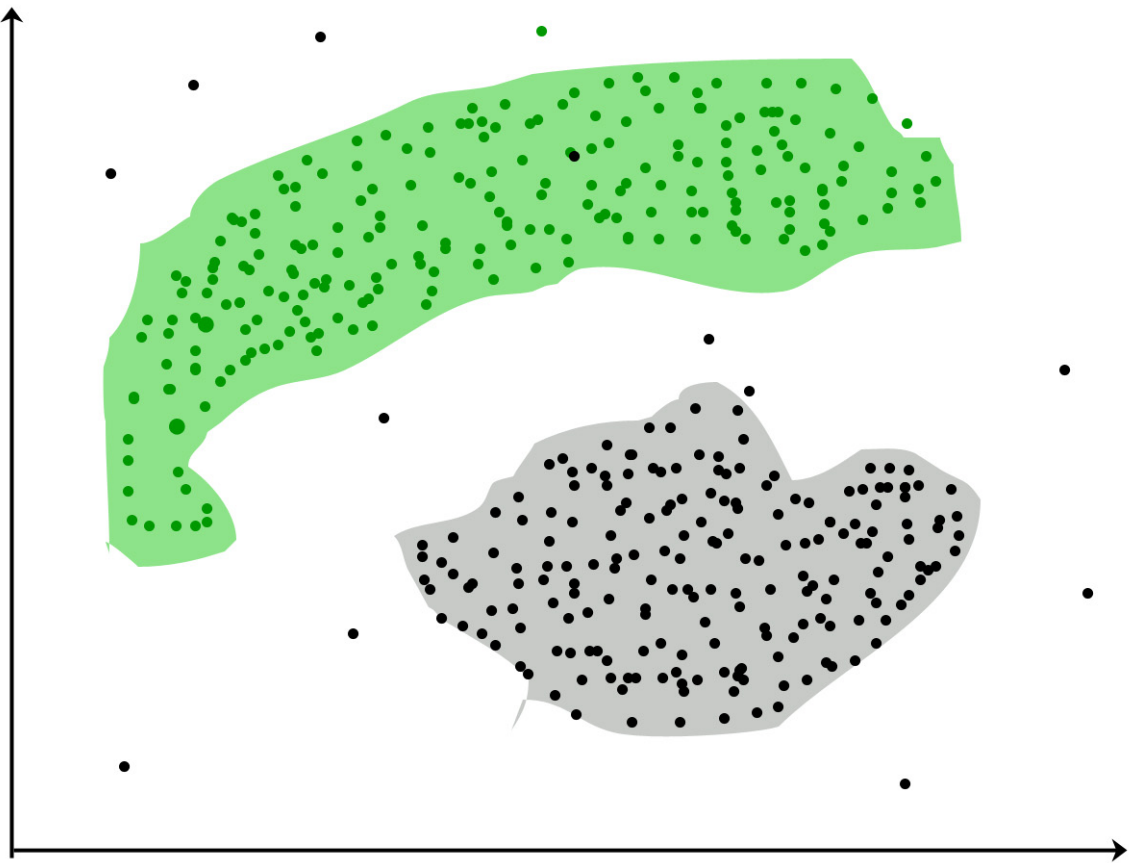
Absence of labels or of formal/descriptive patterns within the data

The model learns these descriptions

These models extract meaningful information or structure from the data, such as:

- automatic feature representation or engineering

- clustering the data into groups



- learning denoising procedures

Original



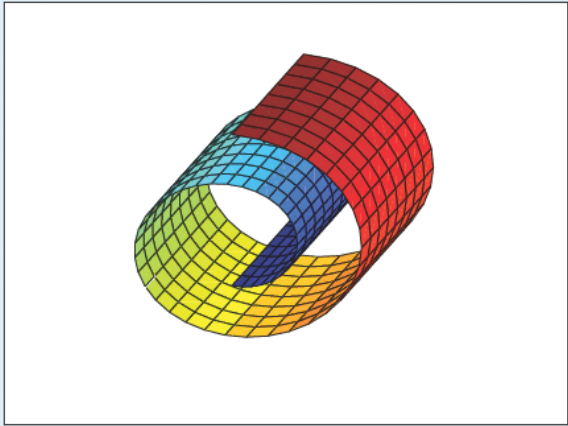
Noisy image



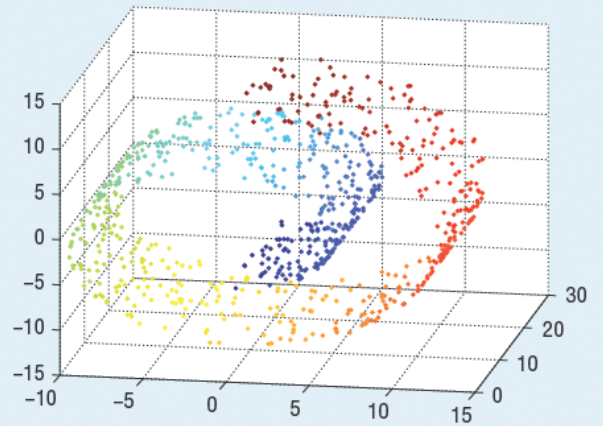
Denoised image



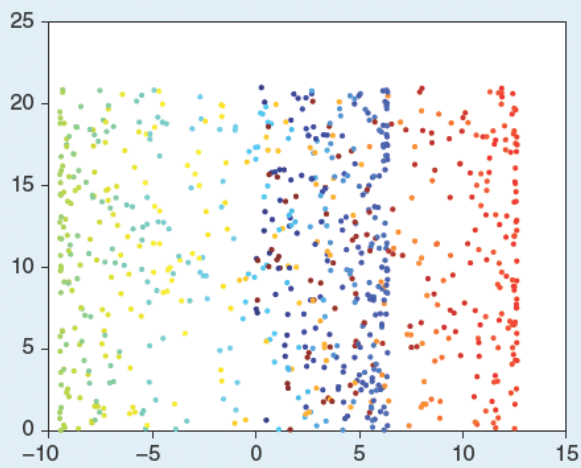
- learning manifolds best describing the data



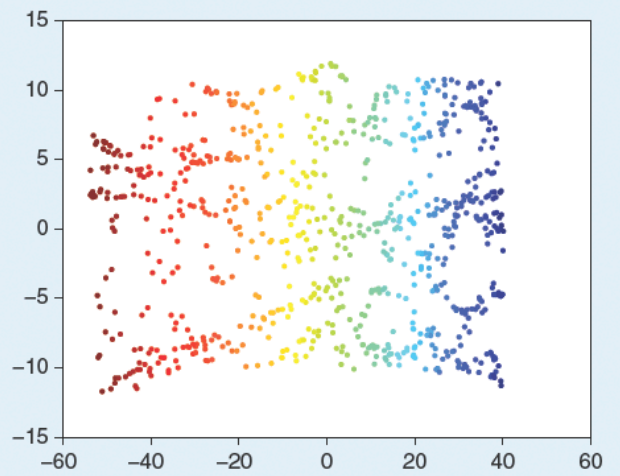
**(a)**



**(b)**

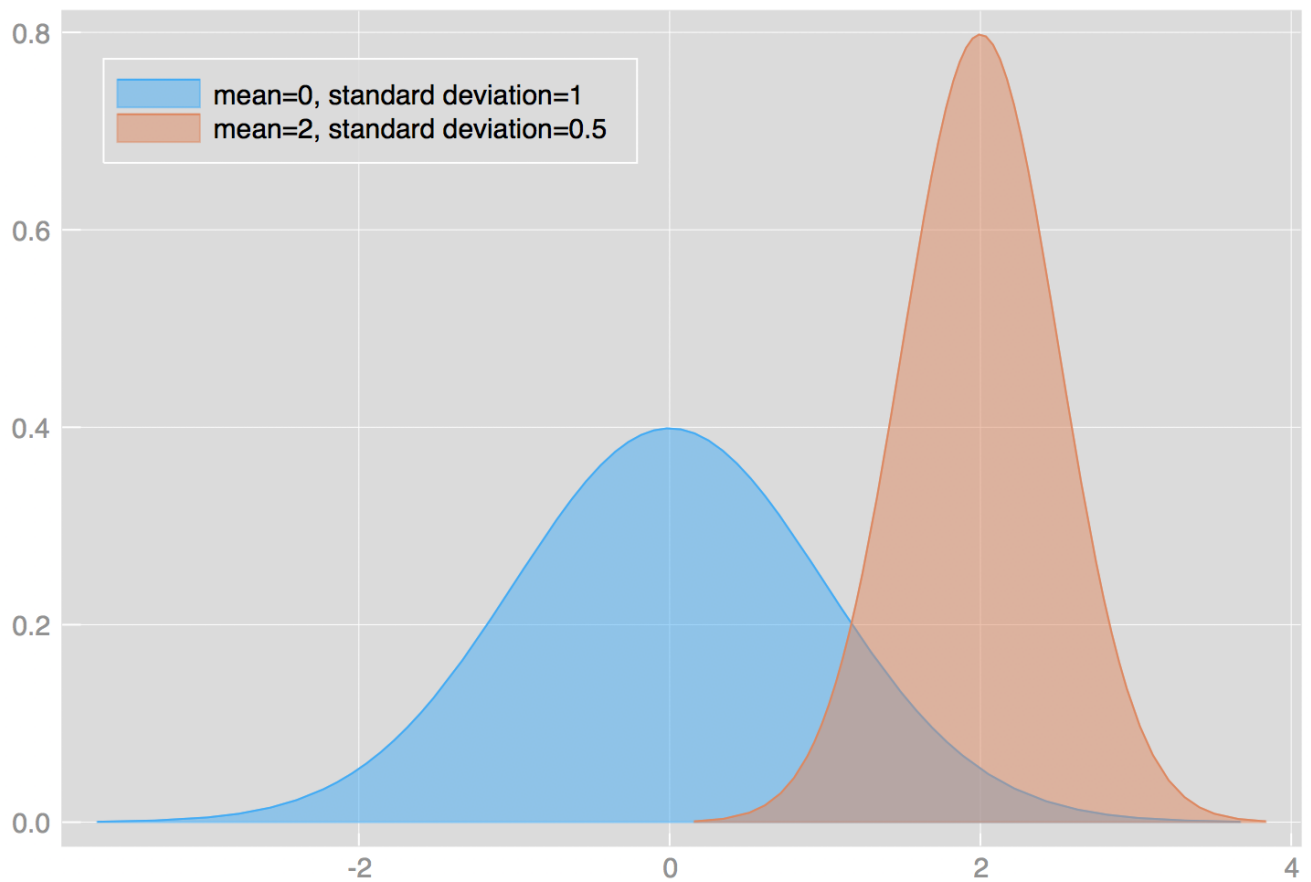


**(c)**



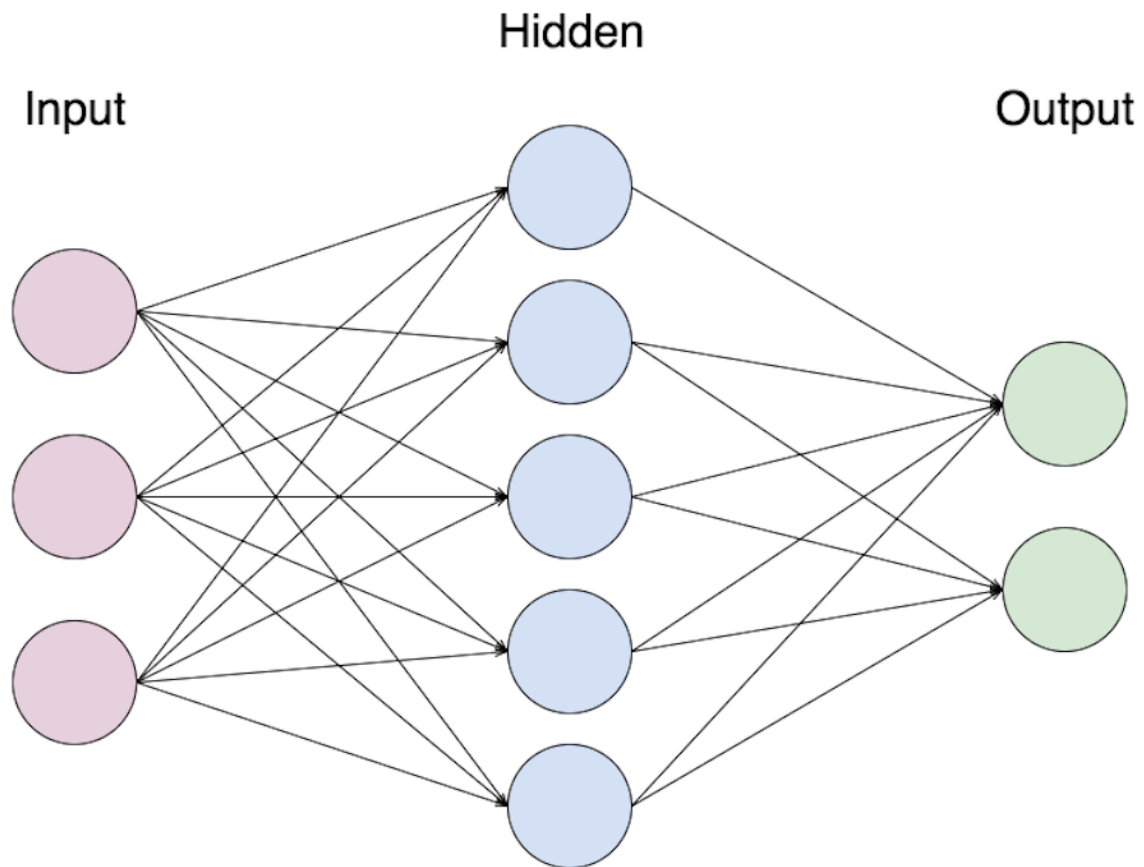
**(d)**

- estimating probability densities



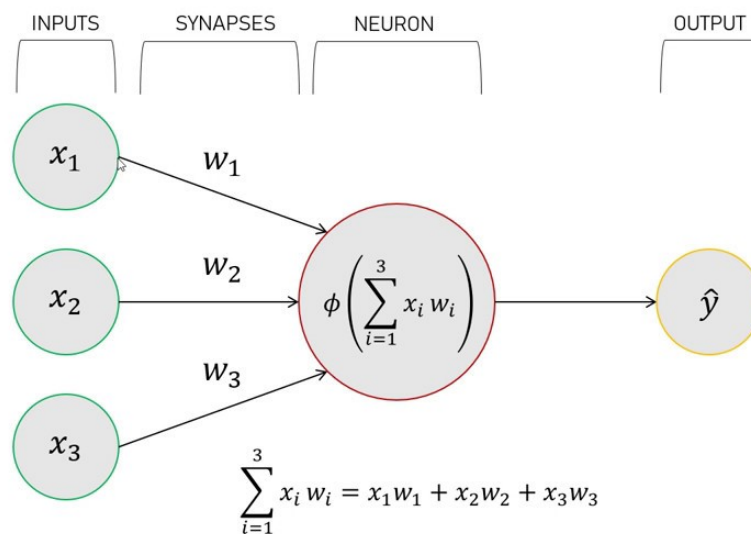
## Standard Neural Network Architecture

1. Input layer
2. Some number of hidden layers
  - each layer has some number of \*nodes/neurons
  - each layer has an activation function
3. Output layer



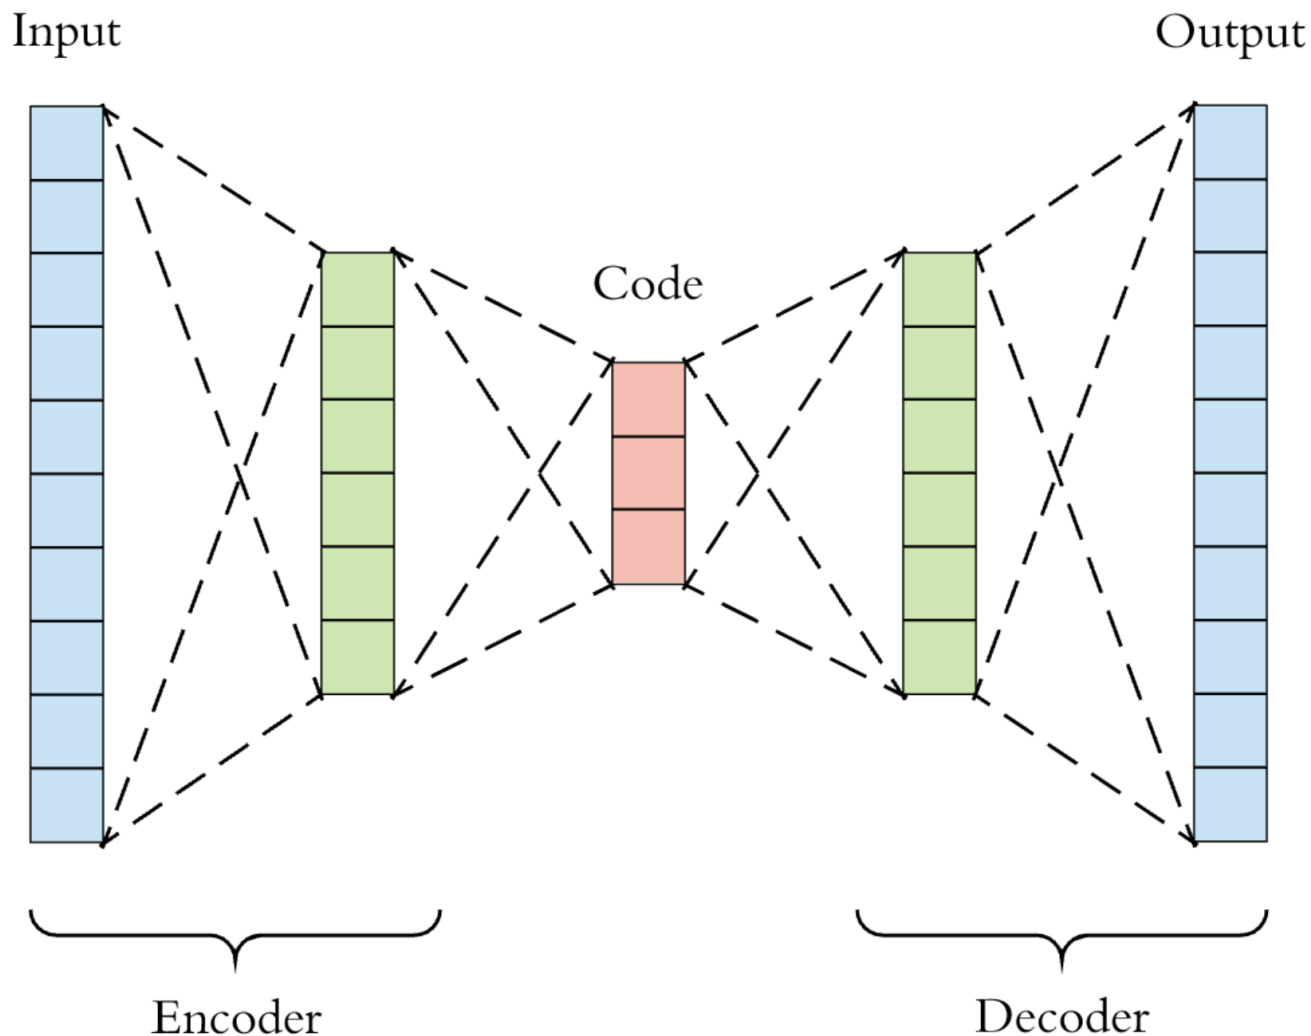
Activation functions,  $\phi$ , can be anything. Generally, these are nonlinear functions such as sigmoid, tanh, ReLU, and ELU.

## PERCEPTRON



- neurons are the neural network's smallest computational unit

# Autoencoders



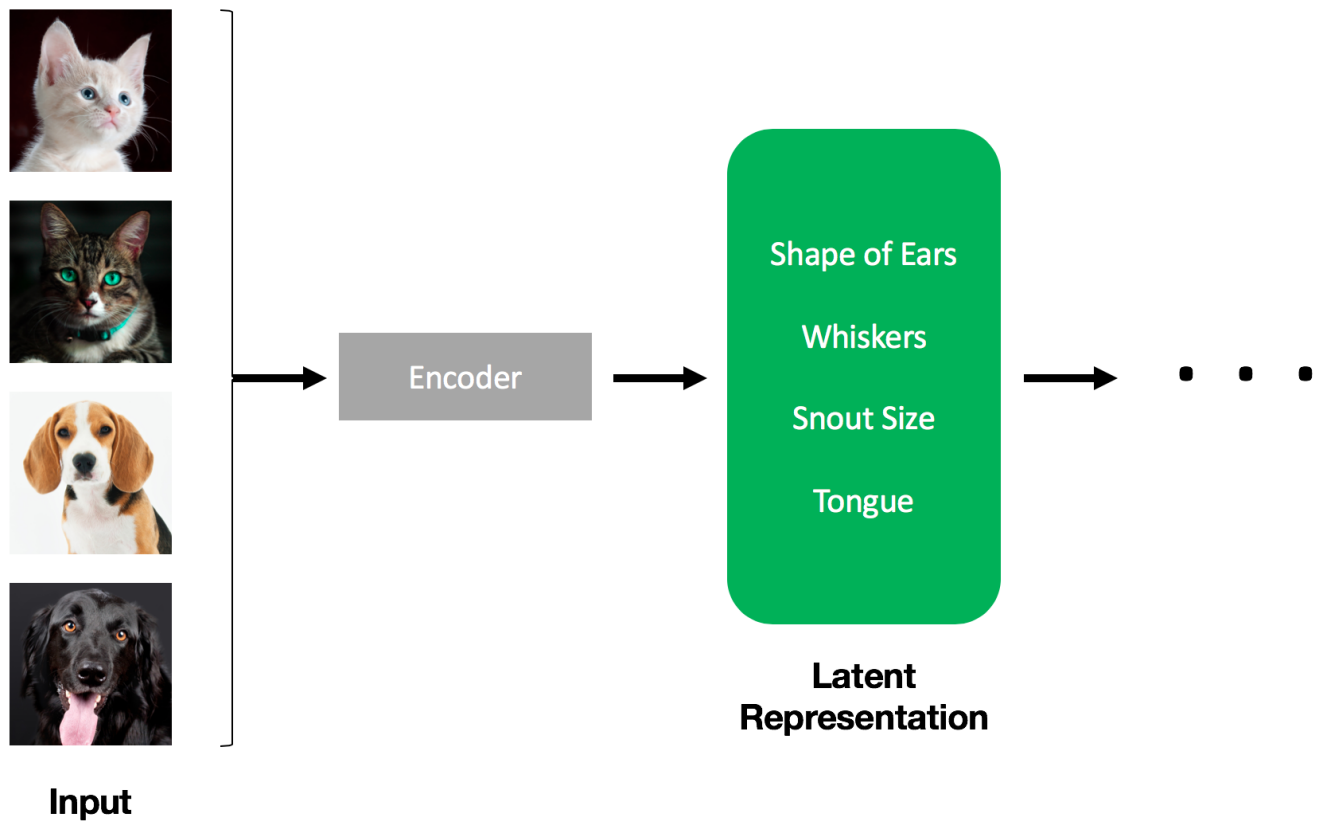
- Neural networks trained to reconstruct the input as the output
- The central hidden layer encodes the input
- The model learns compressed representations of high dimensional data (i.e. representation learning)
- Unsupervised learning of complex distributions
- Major components:
  1. Encoder:  $E(X) : X \rightarrow Z$
  2. Decoder:  $D(Z) : Z \rightarrow X$

$X$  is the original input

$Z$  is the encoded representation (a.k.a. latent representation)
- Auto Encoder:  $f(X) = X$

- Objective function:  $\min ||X - D(Z)||^2$

## Latent Space Representation



- Hidden layer of the autoencoder learns useful properties and potentially prioritizes features

## Example Use Cases for Autoencoders

- Representation learning for input to classification model
- Data compression (i.e. dimensionality reduction)
- Noise removal

### ▼ Code

```
1 !python --version
2
3 # Create folder for saved models
4 !mkdir -p saved_models
```

Python 3.6.9



```

1  %pylab inline
2  font = {'family' : 'normal',
3         #'weight' : 'bold',
4         'size'   : 14}
5  matplotlib.rc('font', **font)
6
7  # Matrix math library
8  import numpy as np
9
10 # Time monitoring and calculation library
11 from time import time
12
13 # Image processing and augmentation library
14 from imgaug import augmenters as iaa
15
16 # Machine learning library for constructing models
17 import tensorflow as tf
18 import tensorflow.keras as keras
19 import tensorflow.keras.backend as K
20 from tensorflow.keras.datasets import fashion_mnist
21 from tensorflow.keras import callbacks, regularizers, Sequential
22 from tensorflow.keras.models import Model
23 from tensorflow.keras.layers import Lambda, Layer, Dense, Input
24 from tensorflow.keras.layers import Conv2D, MaxPool2D, UpSampling2D

```

Populating the interactive namespace from numpy and matplotlib

```

1  # Get Start Time
2  gt0 = time()
3
4  keras.__version__

'2.4.0'

1  # LOAD TRAINING AND VALIDATION DATA
2  # [Fashion MNIST](https://www.tensorflow.org/tutorials/keras/classification)
3  (training_x, training_y), (testing_x, testing_y) = fashion_mnist.load_data()
4
5  # Split Training Set into Training and Validation
6  train_size = 50000
7  train_x = training_x[:train_size]
8  train_y = training_y[:train_size]
9  val_x = training_x[train_size:]
10 val_y = training_y[train_size:]
11 test_y = testing_y
12
13 # BASIC PRE_PROCESS
14 # Scale data to range [0,1]
15 train_x = train_x / 255.
16 val_x = val_x / 255

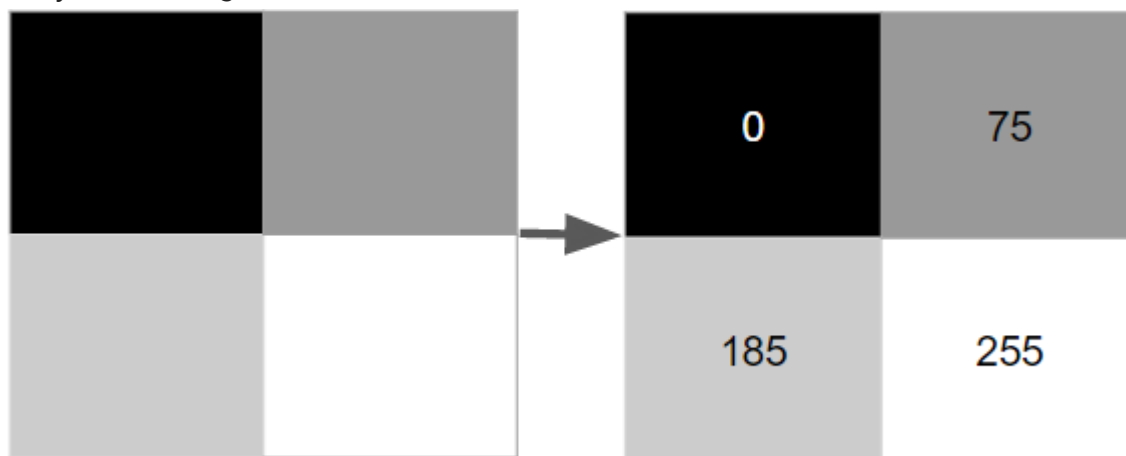
```

```

16     val_x = val_x / 255.
17     test_x = testing_x / 255.

```

## Gray Scale Images



```

1 # Display data shape
2 train_x.shape, train_y.shape, val_x.shape, val_y.shape, test_x.shape, test_y.shape

((50000, 28, 28),
 (50000,),
 (10000, 28, 28),
 (10000,),
 (10000, 28, 28),
 (10000,))

```

For the **training set**, we have 60000 example images that are 28 by 28 pixels.

```
(50000, 28, 28), (50000,)
```

For the **validation set**, we have 10000 examples

```
(10000, 28, 28), (10000,)
```

For the **test set**, we have 10000 examples

```
(10000, 28, 28), (10000,)
```

Each of these images is labeled with a number from 0 to 9 for a different article of clothing (e.g. shirt, sneaker, etc.).

## Sub data sets

- **Training set** used to *build and train* initial models
- **Validation set** used to *select* best of version or configuration of the model

- **Test set** used to verify generalization ability of the "best" model on an *independent* data set  
(NOTE: this data set is not used for training nor selection, hence it's independent of learning the model. Useful to help reduce model bias and increase confidence in model consistency)

```

1 class_names = ['T-shirt/top', 'Trouser', 'Pullover', 'Dress', 'Coat',
2               'Sandal', 'Shirt', 'Sneaker', 'Bag', 'Ankle boot']
3 for i, cls in enumerate(class_names):
4     print(i, cls)

```

```

0 T-shirt/top
1 Trouser
2 Pullover
3 Dress
4 Coat
5 Sandal
6 Shirt
7 Sneaker
8 Bag
9 Ankle boot

```

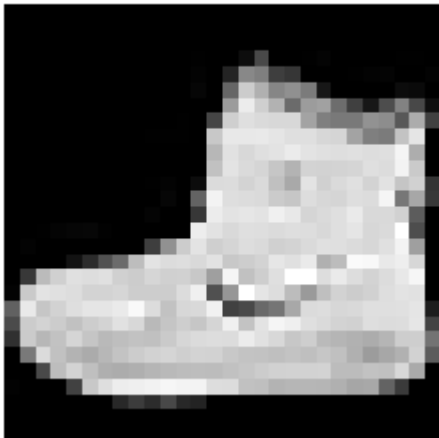
```

1 # Display example image
2 class_number = train_y[0]
3 class_name = class_names[class_number]
4 plt.imshow(train_x[0].reshape(28,28), cmap='gray')
5 plt.title('Class: (%d) %s' % (class_number, class_name))
6 plt.axis('off')

```

```
(-0.5, 27.5, 27.5, -0.5)
```

Class: (9) Ankle boot



```

1 # figure with 11x11 images
2 n = 11
3 img_size = 28
4 # Initialize Grid of Images for Figure
5 figure = np.zeros((img_size * n, img_size * n))
6 # we will sample n points within [-15, 15] standard deviations
7 grid_x = np.linspace(0, 5, n)
8 grid_y = np.linspace(0, 5, n)
9
10

```

```

10 k = 0
11 for i, yi in enumerate(grid_x):
12     for j, xi in enumerate(grid_y):
13         # Reshape and display example image
14         img = train_x[k].reshape(28, 28) # i + j
15         figure[i * img_size: (i + 1) * img_size,
16             j * img_size: (j + 1) * img_size] = img
17         k += 1
18
19 plt.figure(figsize=(10, 10))
20 plt.imshow(figure, cmap='gray')
21 plt.axis("off")
22 plt.show()

```



```

1 # Get the class labels and the corresponding counts
2 classes, class_counts = np.unique(train_y, return_counts=True)
3 classes_v, class_counts_v = np.unique(val_y, return_counts=True)
4
5 # Bar Plot
6 plt.figure(figsize=(10,8), )
7 plt.barh(class_names, class_counts, label='Training')

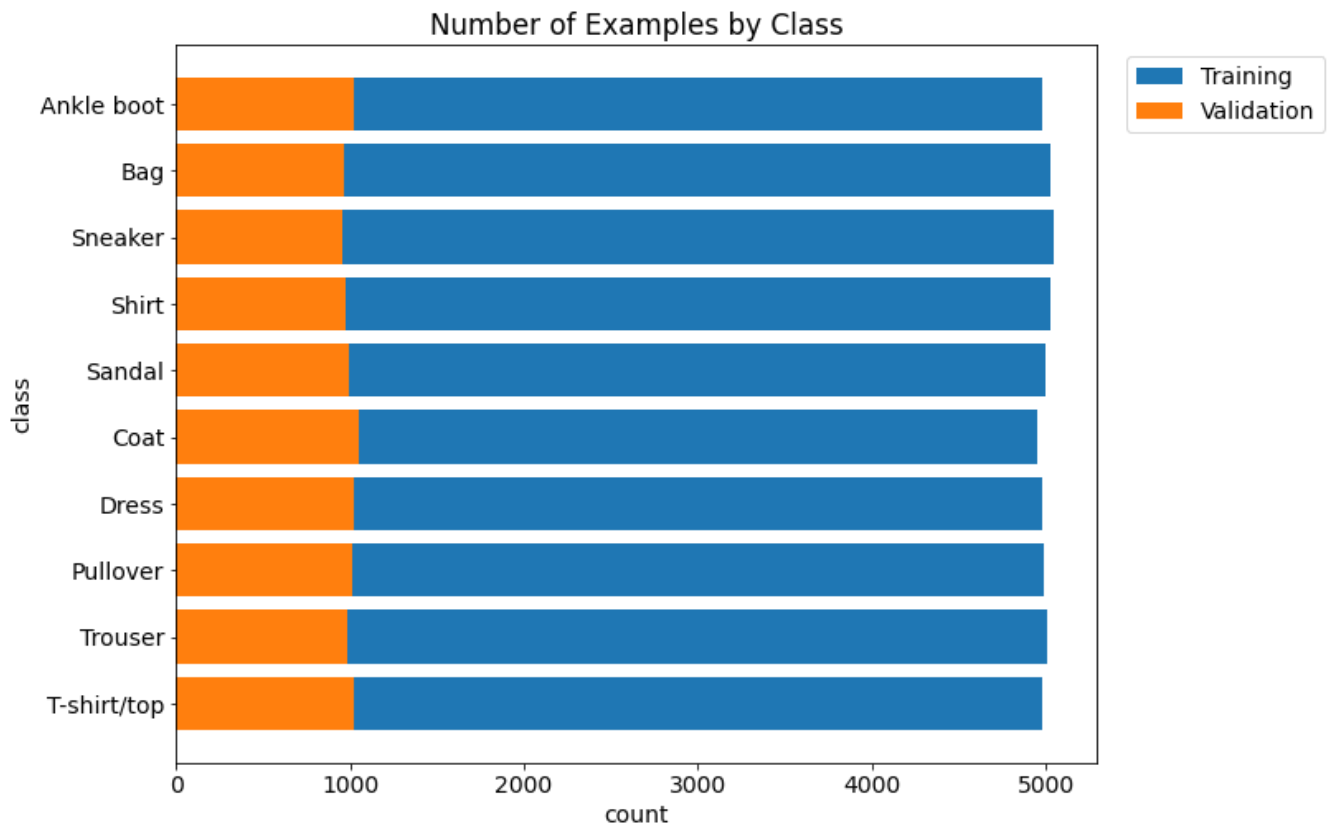
```

```

8
9 plt.barh(class_names, class_counts_v, label='Validation')
10 plt.title('Number of Examples by Class')
11 plt.xlabel('count')
12 plt.ylabel('class')
13
14 plt.legend(bbox_to_anchor=(1.02, 1))

```

<matplotlib.legend.Legend at 0x7f06113d0908>



## ▼ Building Traditional Autoencoder

Now that we know a bit about the dataset and the structure of models, let's build an autoencoder to create a compressed representation of our images.

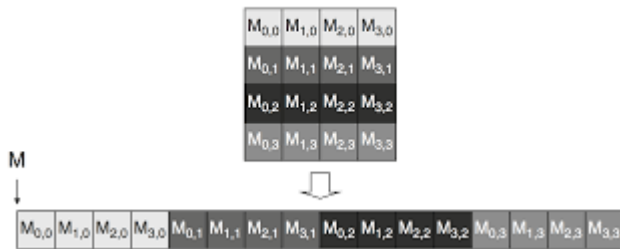
```

1 # RESHAPE DATA INTO VECTOR FORMAT
2 # TODO: The pixels are now rearranged into a 1D vector instead of a 2D matrix
3 train_x = train_x.reshape(-1, 784)
4 val_x = val_x.reshape(-1, 784)
5 test_x = test_x.reshape(-1, 784)
6
7 # Display data shape
8 train_x.shape, val_x.shape, test_x.shape

```

((50000, 784), (10000, 784), (10000, 784))

## Reshaping 2D matrix to 1D vector



```
1  # TODO: Input placeholder
2  input_img = Input(shape=(784,), name='x')
3
4  # Encoded input representation
5  l1_out = Dense(2000, activation='relu', name='encoder_L1')(input_img)
6  l2_out = Dense(500, activation='relu', name='encoder_L2')(l1_out)
7  l3_out = Dense(500, activation='relu', name='encoder_L3')(l2_out)
8  latent = Dense(10, activation='sigmoid', name='z')(l3_out)
9
10 # Model maps input to an encoded representation
11 encoder = Model(input_img, latent)
12
13
14 # Lossy reconstruction of the input
15 l5_out = Dense(500, activation='relu', name='decoder_L1')(latent)
16 l6_out = Dense(500, activation='relu', name='decoder_L2')(l5_out)
17 l7_out = Dense(2000, activation='relu', name='decoder_L3')(l6_out)
18 recon = Dense(784, name='decoder_recon')(l7_out)
19
20
21 # FULL AE MODEL
22 # Model mapping input to its reconstruction
23 autoencoder = Model(input_img, recon)
24
25 # Display summary of model architecture
26 autoencoder.summary()
27
28 # Compile model, specifying training configuration (optimizer, loss, metrics, etc.)
29 autoencoder.compile(optimizer='adam', loss='mse')
```

Model: "functional\_3"

Layer (type)	Output Shape	Param #
=====		
x (InputLayer)	[(None, 784)]	0

encoder_L1 (Dense)	(None, 2000)	1570000
encoder_L2 (Dense)	(None, 500)	1000500
encoder_L3 (Dense)	(None, 500)	250500
z (Dense)	(None, 10)	5010
decoder_L1 (Dense)	(None, 500)	5500
decoder_L2 (Dense)	(None, 500)	250500
decoder_L3 (Dense)	(None, 2000)	1002000
decoder_recon (Dense)	(None, 784)	1568784
=====		
Total params: 5,652,794		
Trainable params: 5,652,794		
Non-trainable params: 0		

```

1  # Keras Callback for early stopping of training
2  estop = keras.callbacks.EarlyStopping(monitor='val_loss', min_delta=0,
3                                         patience=5, verbose=1, mode='auto')
4
5  # TODO: Train the model in "slices" or "batches"
6  # Repeatedly iterate over the entire dataset for a given number of "epochs"
7  t0 = time()
8  train_history = autoencoder.fit(train_x, train_x, epochs=10, batch_size=2048,
9                                  validation_data=(val_x, val_x), callbacks=[estop])
10 t1 = time()
11 etime = (t1 - t0) / 60
12 print("Elapsed time: %.02f min" % etime)

```

Epoch 1/10  
25/25 [=====] - 27s 1s/step - loss: 0.0965 - val\_loss: 0.0730  
Epoch 2/10  
25/25 [=====] - 28s 1s/step - loss: 0.0659 - val\_loss: 0.0553  
Epoch 3/10  
25/25 [=====] - 28s 1s/step - loss: 0.0483 - val\_loss: 0.0424  
Epoch 4/10  
25/25 [=====] - 28s 1s/step - loss: 0.0400 - val\_loss: 0.0380  
Epoch 5/10  
25/25 [=====] - 27s 1s/step - loss: 0.0358 - val\_loss: 0.0335  
Epoch 6/10  
25/25 [=====] - 28s 1s/step - loss: 0.0330 - val\_loss: 0.0314  
Epoch 7/10  
25/25 [=====] - 28s 1s/step - loss: 0.0299 - val\_loss: 0.0291  
Epoch 8/10  
25/25 [=====] - 28s 1s/step - loss: 0.0269 - val\_loss: 0.0255  
Epoch 9/10  
25/25 [=====] - 28s 1s/step - loss: 0.0251 - val\_loss: 0.0243  
Epoch 10/10  
25/25 [=====] - 28s 1s/step - loss: 0.0238 - val\_loss: 0.0234  
Elapsed time: 4.83 min

```

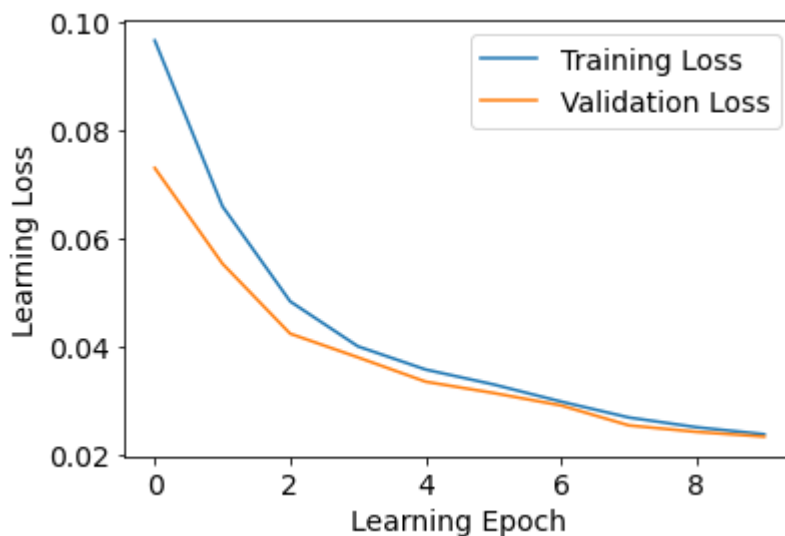
1 # Plot Learning Loss
2 def plot_learning_loss(history):
3     loss = history.history['loss']
4     val_loss = history.history['val_loss']
5
6     plt.plot(loss, label='Training Loss')
7     plt.plot(val_loss, label='Validation Loss')
8     plt.xlabel("Learning Epoch")
9     plt.ylabel("Learning Loss")
10    plt.legend()

```

```

1 plot_learning_loss(train_history)

```



```

1 # RECONSTRUCTION
2 def plot_compare_reconstruction(ae, x):
3     """
4     PARAMS:
5         ae: autoencoder
6         x: example data set
7     """
8     # TODO: 'Predict' the reconstruction, using test set
9     recon = ae.predict(x)
10
11    # Compare original output to reconstructed
12    plt.subplot(2, 2, 1)
13    plt.imshow(recon[0].reshape(28,28), cmap='gray')
14    plt.title('Reconstruction')
15    plt.axis('off')
16    plt.subplot(2, 2, 3)
17    plt.imshow(recon[1].reshape(28,28), cmap='gray')
18    plt.axis('off')
19
20    plt.subplot(2, 2, 2)
21    plt.imshow(x[0].reshape(28,28), cmap='gray')

```



```

21 plt.imshow(x[0].reshape(28,28), cmap='gray')
22 plt.title('Actual')
23 plt.axis('off')
24 plt.subplot(2, 2, 4)
25 plt.imshow(x[1].reshape(28,28), cmap='gray')
26 plt.axis('off')

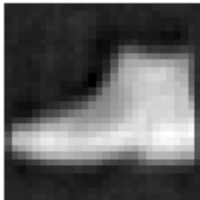
```

```

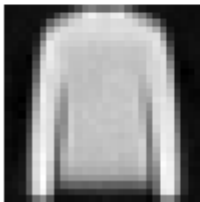
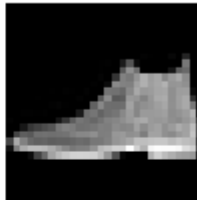
1 # RECONSTRUCTION
2 plot_compare_reconstruction(autoencoder, test_x)

```

Reconstruction



Actual



## ▼ Build Classifier Using Encoding

```

1 # TODO: Encode input
2 train_enc = encoder.predict(train_x)
3 val_enc = encoder.predict(val_x)
4 test_enc = encoder.predict(test_x)
5
6 train_enc.shape, val_enc.shape, test_enc.shape
7
8 ((50000, 10), (10000, 10), (10000, 10))
9
10 # TODO: Input placeholder
11 input_enc = Input(shape=(10,), name='x')
12
13 # Encoded input representation
14 l2_out = Dense(500, activation='relu', name='classifier_L2')(input_enc)
15 l3_out = Dense(250, activation='relu', name='classifier_L3')(l2_out)
16 y = Dense(10, activation='sigmoid', name='z')(l3_out)
17
18 # Model maps input to classification
19 classifier = Model(input_enc, y)
20
21 # Display summary of model architecture
22 classifier.summary()
23
24

```

```

15 # Compile model
16 classifier.compile(optimizer='adam',
17                    loss=[keras.losses.SparseCategoricalCrossentropy()],
18                    metrics=[keras.metrics.SparseCategoricalAccuracy()])

```

Model: "functional\_5"

Layer (type)	Output Shape	Param #
x (InputLayer)	[(None, 10)]	0
classifier_L2 (Dense)	(None, 500)	5500
classifier_L3 (Dense)	(None, 250)	125250
z (Dense)	(None, 10)	2510
Total params: 133,260		
Trainable params: 133,260		
Non-trainable params: 0		

```

1 # Keras Callback for early stopping of training
2 estop = keras.callbacks.EarlyStopping(monitor='val_loss', min_delta=0,
3                                       patience=10, verbose=1, mode='auto')
4
5 # TODO: Train the model
6 t0 = time()
7 train_history = classifier.fit(train_enc, train_y, epochs=20, batch_size=2048,
8                               validation_data=(val_enc, val_y), callbacks=[estop])
9 t1 = time()
10 etime = (t1 - t0) / 60
11 print("Elapsed time: %.02f min" % etime)
12
13 # Get Learning Loss
14 plot_learning_loss(train_history)

```

```

Epoch 1/20
25/25 [=====] - 2s 66ms/step - loss: 2.0720 - sparse_categorical_crossentropy
Epoch 2/20
25/25 [=====] - 1s 43ms/step - loss: 1.2352 - sparse_categorical_crossentropy
Epoch 3/20
25/25 [=====] - 1s 40ms/step - loss: 0.8325 - sparse_categorical_crossentropy
Epoch 4/20
25/25 [=====] - 1s 39ms/step - loss: 0.7438 - sparse_categorical_crossentropy
Epoch 5/20
25/25 [=====] - 1s 40ms/step - loss: 0.7150 - sparse_categorical_crossentropy
Epoch 6/20
25/25 [=====] - 1s 40ms/step - loss: 0.6977 - sparse_categorical_crossentropy
Epoch 7/20
25/25 [=====] - 1s 40ms/step - loss: 0.6855 - sparse_categorical_crossentropy
Epoch 8/20
25/25 [=====] - 1s 43ms/step - loss: 0.6786 - sparse_categorical_crossentropy
Epoch 9/20
25/25 [=====] - 1s 44ms/step - loss: 0.6734 - sparse_categorical_crossentropy
Epoch 10/20
25/25 [=====] - 1s 43ms/step - loss: 0.6627 - sparse_categorical_crossentropy
Epoch 11/20
25/25 [=====] - 1s 43ms/step - loss: 0.6563 - sparse_categorical_crossentropy
Epoch 12/20
25/25 [=====] - 1s 45ms/step - loss: 0.6478 - sparse_categorical_crossentropy
Epoch 13/20
25/25 [=====] - 1s 43ms/step - loss: 0.6445 - sparse_categorical_crossentropy
Epoch 14/20
25/25 [=====] - 1s 41ms/step - loss: 0.6456 - sparse_categorical_crossentropy
Epoch 15/20
25/25 [=====] - 1s 41ms/step - loss: 0.6390 - sparse_categorical_crossentropy
Epoch 16/20
25/25 [=====] - 1s 42ms/step - loss: 0.6320 - sparse_categorical_crossentropy
Epoch 17/20
25/25 [=====] - 1s 43ms/step - loss: 0.6280 - sparse_categorical_crossentropy
Epoch 18/20
25/25 [=====] - 1s 43ms/step - loss: 0.6244 - sparse_categorical_crossentropy
Epoch 19/20
25/25 [=====] - 1s 43ms/step - loss: 0.6194 - sparse_categorical_crossentropy
Epoch 20/20
25/25 [=====] - 1s 43ms/step - loss: 0.6191 - sparse_categorical_crossentropy
Elapsed time: 0.38 min

```

```

1 # Evaluate classifier using test data
2 #test_encoding = encoder.predict(test_x)
3 test_loss, test_acc = classifier.evaluate(test_enc, test_y, batch_size=128)
4
5 print("Test Loss: %.04f \nTest Accuracy: %.02f%%" % (test_loss, test_acc * 100))

```

```

79/79 [=====] - 0s 2ms/step - loss: 0.6395 - sparse_categorical_crossentropy
Test Loss: 0.6395
Test Accuracy: 74.88%

```



## ▼ Build Classifier Directly from Image Vector

Let's quickly compare and see the performance of a model trained using the encodings versus the performance of a model trained using all the pixels

```
1  # Input placeholder
2  input_x = Input(shape=(784,), name='x')
3
4  # Encoded input representation
5  l2_out = Dense(500, activation='relu', name='classifier_L2')(input_x)
6  l3_out = Dense(250, activation='relu', name='classifier_L3')(l2_out)
7  y = Dense(10, activation='sigmoid', name='z')(l3_out)
8
9  # Model maps input to classification
10 img_classifier = Model(input_x, y)
11
12 # Display summary of model architecture
13 img_classifier.summary()
14
15 # Compile model
16 img_classifier.compile(optimizer='adam',
17                       loss=[keras.losses.SparseCategoricalCrossentropy()],
18                       metrics=[keras.metrics.SparseCategoricalAccuracy()])
19
20 # Keras Callback for early stopping of training
21 estop = keras.callbacks.EarlyStopping(monitor='val_loss', min_delta=0,
22                                       patience=5, verbose=1, mode='auto')
23
24 # Train the model
25 t0 = time()
26 train_history = img_classifier.fit(train_x, train_y, epochs=10, batch_size=2048,
27                                   validation_data=(val_x, val_y), callbacks=[estop])
28 t1 = time()
29 etime = (t1 - t0) / 60
30 print("Elapsed time: %.02f min" % etime)
31
32 # Plot Learning Loss
33 plot_learning_loss(train_history)
```

Model: "functional\_7"

Layer (type)	Output Shape	Param #
=====		
x (InputLayer)	[(None, 784)]	0
=====		
classifier_L2 (Dense)	(None, 500)	392500
=====		
classifier_L3 (Dense)	(None, 250)	125250
=====		
z (Dense)	(None, 10)	2510
=====		
Total params: 520,260		
Trainable params: 520,260		
Non-trainable params: 0		

```
Epoch 1/10
25/25 [=====] - 3s 109ms/step - loss: 1.0266 - sparse_categorical_crossentropy: 1.0266
Epoch 2/10
25/25 [=====] - 3s 105ms/step - loss: 0.5259 - sparse_categorical_crossentropy: 0.5259
Epoch 3/10
25/25 [=====] - 3s 105ms/step - loss: 0.4456 - sparse_categorical_crossentropy: 0.4456
Epoch 4/10
25/25 [=====] - 3s 105ms/step - loss: 0.4024 - sparse_categorical_crossentropy: 0.4024
Epoch 5/10
25/25 [=====] - 3s 105ms/step - loss: 0.3864 - sparse_categorical_crossentropy: 0.3864
Epoch 6/10
25/25 [=====] - 3s 105ms/step - loss: 0.3578 - sparse_categorical_crossentropy: 0.3578
Epoch 7/10
25/25 [=====] - 3s 104ms/step - loss: 0.3507 - sparse_categorical_crossentropy: 0.3507
Epoch 8/10
25/25 [=====] - 3s 105ms/step - loss: 0.3312 - sparse_categorical_crossentropy: 0.3312
Epoch 9/10
25/25 [=====] - 3s 105ms/step - loss: 0.3181 - sparse_categorical_crossentropy: 0.3181
Epoch 10/10
25/25 [=====] - 3s 102ms/step - loss: 0.3039 - sparse_categorical_crossentropy: 0.3039
Elapsed time: 0.46 min
```



```
1 # Evaluate classifier using test data
2 test_loss, test_acc = img_classifier.evaluate(test_x, test_y, batch_size=128)
3 print("Test Loss: %.04f \nTest Accuracy: %.02f%%" % (test_loss, test_acc * 100))

79/79 [=====] - 0s 6ms/step - loss: 0.3641 - sparse_categorical_crossentropy: 0.3641
Test Loss: 0.3641
Test Accuracy: 87.45%
```



## ▼ Sparse Autoencoder

Optimize compressed output by reducing the amount of memory using sparse representations instead

To make representations more compact, impose a sparsity constraint on the activation of the hidden representations (this is the activity regularizer in Keras), such that fewer units get activated at a given time

In Keras we use the "activity\_regularizer" parameter for each layer to apply penalties on parameters or activations during optimization. Penalties are incorporated in the loss function

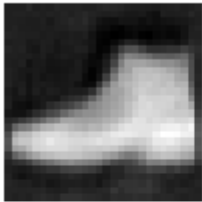
```
1 # Load existing model
2 sparse_ae = tf.keras.models.load_model('saved_models/sparse_autoencoder')
3 sparse_ae.summary()
4 plot_compare_reconstruction(sparse_ae, test_x)
```

Model: "functional\_11"

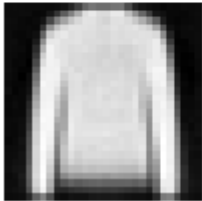
Layer (type)	Output Shape	Param #
=====		
x (InputLayer)	[(None, 784)]	0
encoder_L1 (Dense)	(None, 2000)	1570000
encoder_L2 (Dense)	(None, 500)	1000500
encoder_L3 (Dense)	(None, 500)	250500
z (Dense)	(None, 10)	5010
decoder_L1 (Dense)	(None, 500)	5500
decoder_L2 (Dense)	(None, 500)	250500
decoder_L3 (Dense)	(None, 2000)	1002000
sparse_recon (Dense)	(None, 784)	1568784
=====		

Total params: 5,652,794  
Trainable params: 5,652,794  
Non-trainable params: 0

Reconstruction



Actual



```

1  # Include activity constraint by defining a small value for the activity_regularizer
2  # ref (https://www.tensorflow.org/api\_docs/python/tf/keras/regularizers/Regularizer)
3
4  # Input placeholder
5  input_img = Input(shape=(784,), name='x')
6
7  # TODO: Encoded input representation
8  l1_out = Dense(2000, activation='relu', name='encoder_L1')(input_img)
9  l2_out = Dense(500, activation='relu', name='encoder_L2',
10               activity_regularizer=regularizers.l1(10e-10))(l1_out)
11  l3_out = Dense(500, activation='relu', name='encoder_L3',
12               activity_regularizer=regularizers.l1(10e-10))(l2_out)
13  latent = Dense(10, activation='sigmoid', name='z',
14               activity_regularizer=regularizers.l1(10e-10))(l3_out)
15
16  # Model maps input to an encoded representation
17  sparse_encoder = Model(input_img, latent)
18
19
20  # Lossy reconstruction of the input
21  l5_out = Dense(500, activation='relu', name='decoder_L1')(latent)
22  l6_out = Dense(500, activation='relu', name='decoder_L2')(l5_out)
23  l7_out = Dense(2000, activation='relu', name='decoder_L3')(l6_out)
24  sparse_recon = Dense(784, name='sparse_recon')(l7_out)
25
26  # Model mapping latent representation to input reconstruction
27  #sparse_decoder = Model(encoded, decoded)
28
29  # Model mapping input to its reconstruction
30  sparse_autoencoder = Model(input_img, sparse_recon)
31
32  # Display summary of model architecture
33  sparse_autoencoder.summary()
34
35  # Compile model
36  sparse_autoencoder.compile(optimizer='adam', loss='mse')

```

Model: "functional\_11"

Layer (type)	Output Shape	Param #
=====		
x (InputLayer)	[(None, 784)]	0
encoder_L1 (Dense)	(None, 2000)	1570000
encoder_L2 (Dense)	(None, 500)	1000500
encoder_L3 (Dense)	(None, 500)	250500
z (Dense)	(None, 10)	5010
decoder_L1 (Dense)	(None, 500)	5500

decoder_L2 (Dense)	(None, 500)	250500
decoder_L3 (Dense)	(None, 2000)	1002000
sparse_recon (Dense)	(None, 784)	1568784
=====		
Total params: 5,652,794		
Trainable params: 5,652,794		
Non-trainable params: 0		

```

1  # Train the model
2  t0 = time()
3  train_history = sparse_autoencoder.fit(train_x, train_x,
4                                         epochs=10,
5                                         batch_size=2048,
6                                         validation_data=(val_x, val_x))
7  t1 = time()
8  duration = (t1 - t0) / 60 # convert to minutes
9  print('Elapsed Time: %.02f min' % duration)
10
11 # Save the entire model
12 sparse_autoencoder.save('saved_models/sparse_autoencoder')
13
14 # Plot Learning Loss
15 plot_learning_loss(train_history)

```



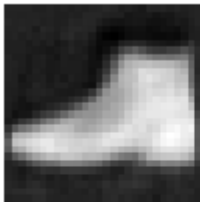
```

Epoch 1/10
25/25 [=====] - 28s 1s/step - loss: 0.0941 - val_loss: 0.0711
Epoch 2/10
25/25 [=====] - 28s 1s/step - loss: 0.0626 - val_loss: 0.0520
Epoch 3/10
25/25 [=====] - 28s 1s/step - loss: 0.0453 - val_loss: 0.0407
Epoch 4/10
25/25 [=====] - 28s 1s/step - loss: 0.0383 - val_loss: 0.0388
Epoch 5/10
25/25 [=====] - 28s 1s/step - loss: 0.0348 - val_loss: 0.0329
Epoch 6/10
25/25 [=====] - 28s 1s/step - loss: 0.0319 - val_loss: 0.0307
Epoch 7/10
25/25 [=====] - 28s 1s/step - loss: 0.0295 - val_loss: 0.0290
Epoch 8/10
25/25 [=====] - 28s 1s/step - loss: 0.0269 - val_loss: 0.0258
Epoch 9/10
25/25 [=====] - 28s 1s/step - loss: 0.0254 - val_loss: 0.0247
Epoch 10/10

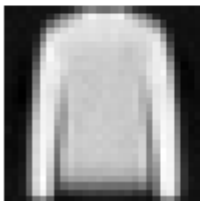
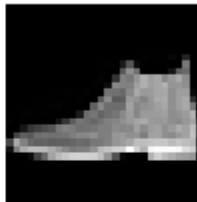
```

```
1 plot_compare_reconstruction(sparse_autoencoder, test_x)
```

Reconstruction



Actual



0

|



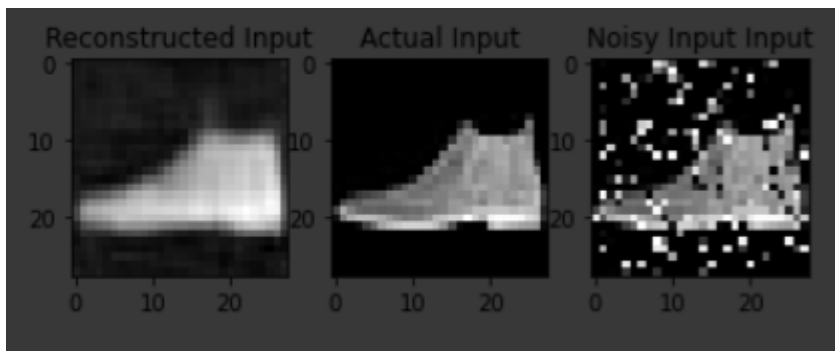
|

## ▼ Denoising AutoEncoders

When an image gets corrupted, or contains noise, there is no straight-forward way to remove the noise.

We want to "denoise" the image and convert the noisy image into a somewhat clearer image with most (or all) of the noise removed.

Example from a simple model:



```

1 # Load existing model
2 denoise_ae = tf.keras.models.load_model('saved_models/denoise_autoencoder')
3 denoise_ae.summary()

```

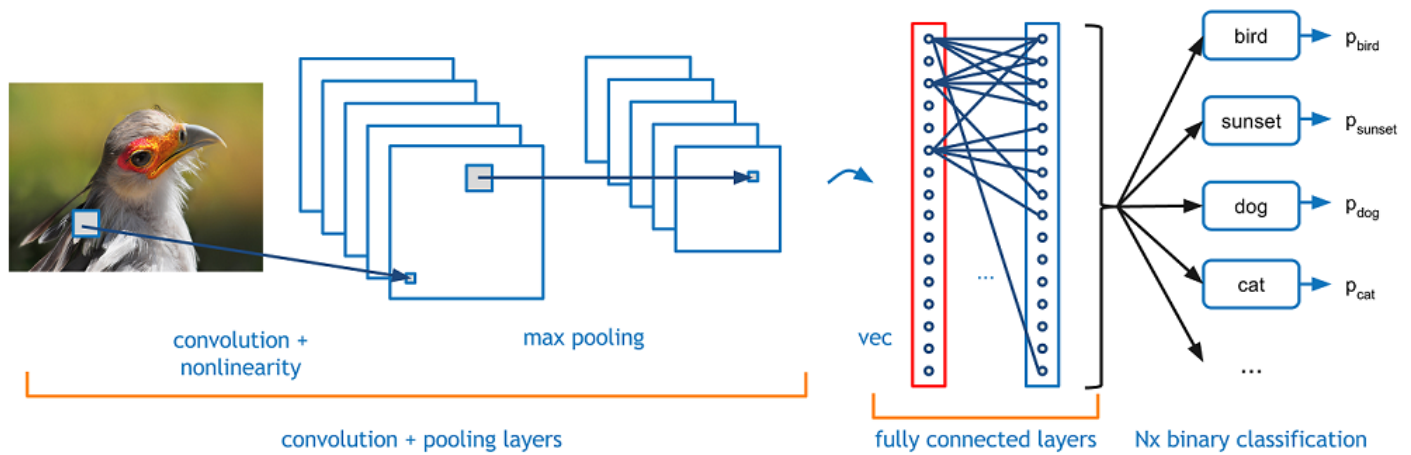
Model: "functional\_15"

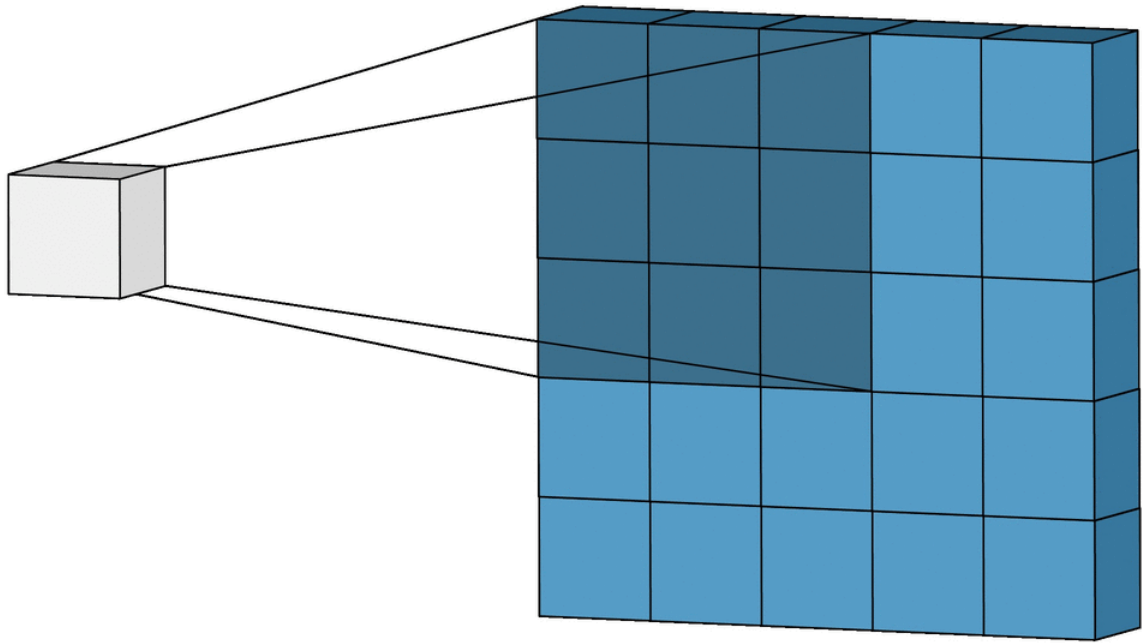
Layer (type)	Output Shape	Param #
=====		
input_1 (InputLayer)	[(None, 28, 28, 1)]	0
encoder_L1_conv2d (Conv2D)	(None, 28, 28, 64)	640
encoder_L1_max (MaxPooling2D)	(None, 14, 14, 64)	0
encoder_L2_conv2d (Conv2D)	(None, 14, 14, 32)	18464
encoder_L2_max (MaxPooling2D)	(None, 7, 7, 32)	0
encoder_L3_conv2d (Conv2D)	(None, 7, 7, 16)	4624
encoder_L3_max (MaxPooling2D)	(None, 4, 4, 16)	0
decoder_L1_conv2d (Conv2D)	(None, 4, 4, 16)	2320
decoder_L1_up (UpSampling2D)	(None, 8, 8, 16)	0
decoder_L2_conv2d (Conv2D)	(None, 8, 8, 32)	4640
decoder_L2_up (UpSampling2D)	(None, 16, 16, 32)	0
decoder_L3_conv2d (Conv2D)	(None, 14, 14, 64)	18496
decoder_L3_up (UpSampling2D)	(None, 28, 28, 64)	0
decoder_recon (Conv2D)	(None, 28, 28, 1)	577
=====		
Total params: 49,761		
Trainable params: 49,761		
Non-trainable params: 0		

## ▼ Convolutional Neural Network (CNN)

Similar structure to a standard neural network

1. Input layer
2. Some number of hidden layers
  - convolution
  - max or mean pooling
  - activation function
  - fully connected dense layers
3. Output layer





```
1 # LOAD DATA
2 (training_x, training_y), (testing_x, testing_y) = fashion_mnist.load_data()
3
4 # Introduce noise to some of the data
5 # NOTE: In reality we don't know the source or structure of the noise
6 seq = iaa.Sequential([iaa.SaltAndPepper(.2)])
7
8 training_x_aug = seq.augment_images(training_x)
9 test_x_aug = seq.augment_images(testing_x)
10
11 # Split Training Data into Training and Validation
12 train_size = 50000
13 train_x_aug = training_x_aug[:train_size]
14 val_x_aug = training_x_aug[train_size:]
15
16 # Clean data
17 train_x = training_x[:train_size]
18 train_y = training_y[:train_size]
19 val_x = training_x[train_size:]
20 val_y = training_y[train_size:]
```

```

1  # PRE_PROCESS THE DATA
2  train_x_aug = train_x_aug / 255.
3  val_x_aug = val_x_aug / 255.
4  test_x_aug = test_x_aug / 255.
5
6  train_x = train_x / 255.
7  val_x = val_x / 255.
8  test_x = testing_x / 255.
9
10 # RESHAPE INTO TENSORS FOR CNN
11 train_x_aug_img = train_x_aug.reshape(-1, 28, 28, 1)
12 val_x_aug_img = val_x_aug.reshape(-1, 28, 28, 1)
13 test_x_aug_img = test_x_aug.reshape(-1, 28, 28, 1)
14
15 train_x_img = train_x.reshape(-1, 28, 28, 1) # TODO
16 val_x_img = val_x.reshape(-1, 28, 28, 1)
17 test_x_img = test_x.reshape(-1, 28, 28, 1)
18
19 train_x_img.shape, val_x_img.shape, test_x_img.shape

((50000, 28, 28, 1), (10000, 28, 28, 1), (10000, 28, 28, 1))

1  # TODO: Input placeholder
2  input_img = Input(shape=(28, 28, 1))
3
4  # Encoded input representation
5  # padding=same: zero padding during convolution and pooling
6  # padding=valid: no padding during convolution and pooling
7  l1_out = Conv2D(64, (3, 3), activation='relu', padding='same',
8              name='encoder_L1_conv2d')(input_img)
9  l1_out = MaxPool2D((2, 2), padding='same', name='encoder_L1_max')(l1_out)
10
11 l2_out = Conv2D(32, (3, 3), activation='relu', padding='same',
12              name='encoder_L2_conv2d')(l1_out)
13 l2_out = MaxPool2D((2, 2), padding='same', name='encoder_L2_max')(l2_out)
14
15 l3_out = Conv2D(16, (3, 3), activation='relu', padding='same',
16              name='encoder_L3_conv2d')(l2_out)
17 l3_out = MaxPool2D((2, 2), padding='same', name='encoder_L3_max')(l3_out)
18
19 # Model mapping input to its encoded representation
20 denoise_encoder = Model(input_img, l3_out)
21
22
23 # Lossy reconstruction of the input
24 l4_out = Conv2D(16, (3, 3), activation='relu', padding='same',
25              name='decoder_L1_conv2d')(l3_out)
26 l4_out = UpSampling2D((2, 2), name='decoder_L1_up')(l4_out)
27
28 l5_out = Conv2D(32, (3, 3), activation='relu', padding='same',
29              name='decoder_L2_conv2d')(l4_out)

```

```

30 15_out = UpSampling2D((2, 2), name='decoder_L2_up')(15_out)
31
32 16_out = Conv2D(64, (3, 3), activation='relu',
33               name='decoder_L3_conv2d')(15_out)
34 16_out = UpSampling2D((2, 2), name='decoder_L3_up')(16_out)
35
36 recon = Conv2D(1, (3, 3), padding='same', name='decoder_recon')(16_out)
37
38
39 # FULL AUTOENCODER
40 # Model mapping input to its reconstruction
41 denoise_autoencoder = Model(input_img, recon)
42
43 # Display Summary
44 denoise_autoencoder.summary()
45
46 # Compile model
47 denoise_autoencoder.compile(optimizer='adam', loss='mse')

```

Model: "functional\_15"

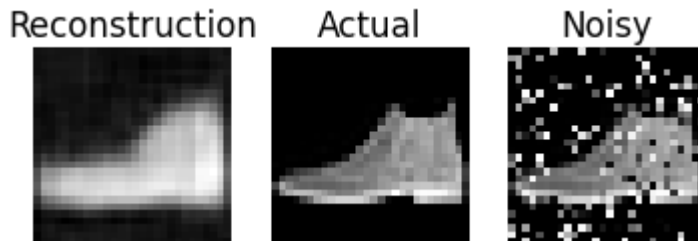
Layer (type)	Output Shape	Param #
=====		
input_1 (InputLayer)	[(None, 28, 28, 1)]	0
-----		
encoder_L1_conv2d (Conv2D)	(None, 28, 28, 64)	640
-----		
encoder_L1_max (MaxPooling2D)	(None, 14, 14, 64)	0
-----		
encoder_L2_conv2d (Conv2D)	(None, 14, 14, 32)	18464
-----		
encoder_L2_max (MaxPooling2D)	(None, 7, 7, 32)	0
-----		
encoder_L3_conv2d (Conv2D)	(None, 7, 7, 16)	4624
-----		
encoder_L3_max (MaxPooling2D)	(None, 4, 4, 16)	0
-----		
decoder_L1_conv2d (Conv2D)	(None, 4, 4, 16)	2320
-----		
decoder_L1_up (UpSampling2D)	(None, 8, 8, 16)	0
-----		
decoder_L2_conv2d (Conv2D)	(None, 8, 8, 32)	4640
-----		
decoder_L2_up (UpSampling2D)	(None, 16, 16, 32)	0
-----		
decoder_L3_conv2d (Conv2D)	(None, 14, 14, 64)	18496
-----		
decoder_L3_up (UpSampling2D)	(None, 28, 28, 64)	0
-----		
decoder_recon (Conv2D)	(None, 28, 28, 1)	577
=====		
Total params: 49,761		
Trainable params: 49,761		
Non-trainable params: 0		
-----		

```
1  # Train the model
2  # NOTE: Noisy image is the input; try to reconstruct the original denoised image
3  # Select subset of data to speed up learning
4  x_in = train_x_aug_img[:20000]
5  x_out = train_x_img[:20000]
6  val_x_in = val_x_aug_img[:5000]
7  val_x_out = val_x_img[:5000]
8
9  t0 = time()
10 train_history = denoise_autoencoder.fit(x_in, x_out, epochs=10, batch_size=1024,
11                                       validation_data=(val_x_aug_img, val_x_img))
12 t1 = time()
13 etime = (t1 - t0) / 60
14 print("Elapsed time: %.02f" % etime)
15
16 # Save the entire model
17 denoise_autoencoder.save('saved_models/denoise_autoencoder')
18
19 # Plot Learning Loss
20 plot_learning_loss(train_history)
```

```
Epoch 1/10
20/20 [=====] - 88s 4s/step - loss: 0.1106 - val_loss: 0.0712
Epoch 2/10
20/20 [=====] - 88s 4s/step - loss: 0.0538 - val_loss: 0.0428
Epoch 3/10
20/20 [=====] - 88s 4s/step - loss: 0.0387 - val_loss: 0.0254
```

```
1 # RECONSTRUCTION
2 denoised_x = denoise_autoencoder.predict(test_x_aug_img)
3
4 # Compare reconstruction to the original input and the noisy input
5 plt.subplot(1, 3, 1)
6 plt.imshow(denoised_x[0].reshape(28, 28), cmap='gray')
7 plt.title('Reconstruction')
8 plt.axis('off')
9
10 plt.subplot(1, 3, 2)
11 plt.imshow(test_x[0].reshape(28, 28), cmap='gray')
12 plt.title('Actual')
13 plt.axis('off')
14
15 plt.subplot(1, 3, 3)
16 plt.imshow(test_x_aug[0].reshape(28, 28), cmap='gray')
17 plt.title('Noisy')
18 plt.axis('off')
```

(-0.5, 27.5, 27.5, -0.5)



Learning Epoch

## ▼ [Variational AutoEncoders \(VAE\)](#)

### Image Generation with Variational AutoEncoders

Main distinction between traditional autoencoders and variational autoencoders is that instead of a compressed bottleneck of information, we attempt to model the probability distribution of the training data.

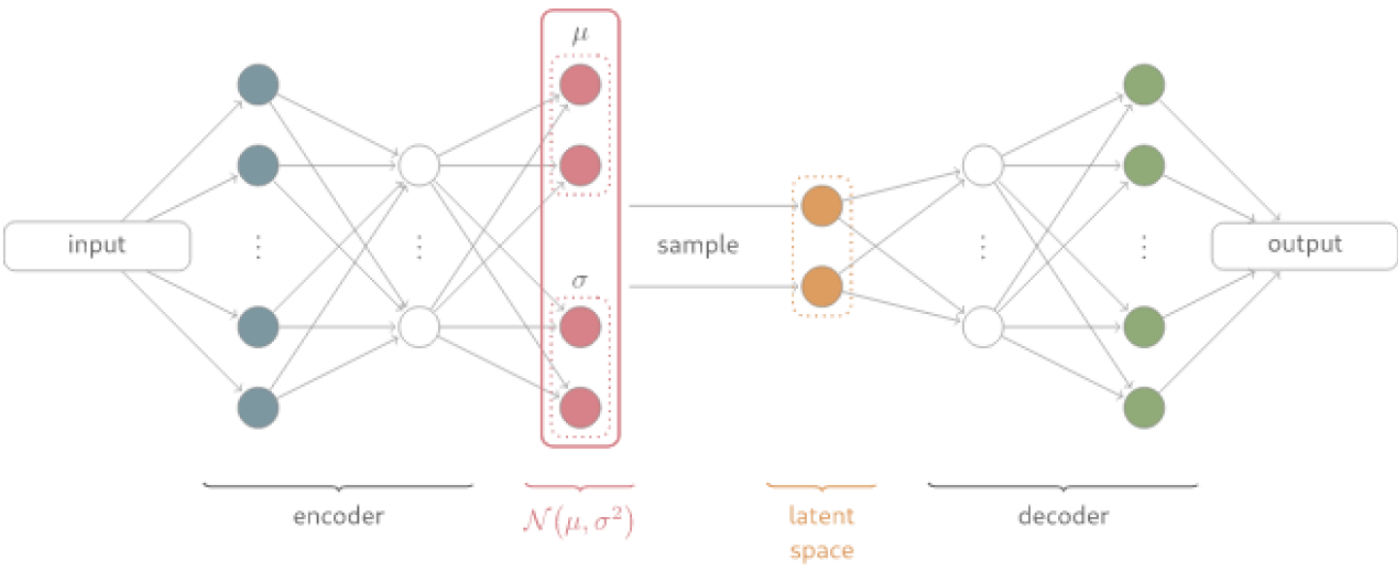
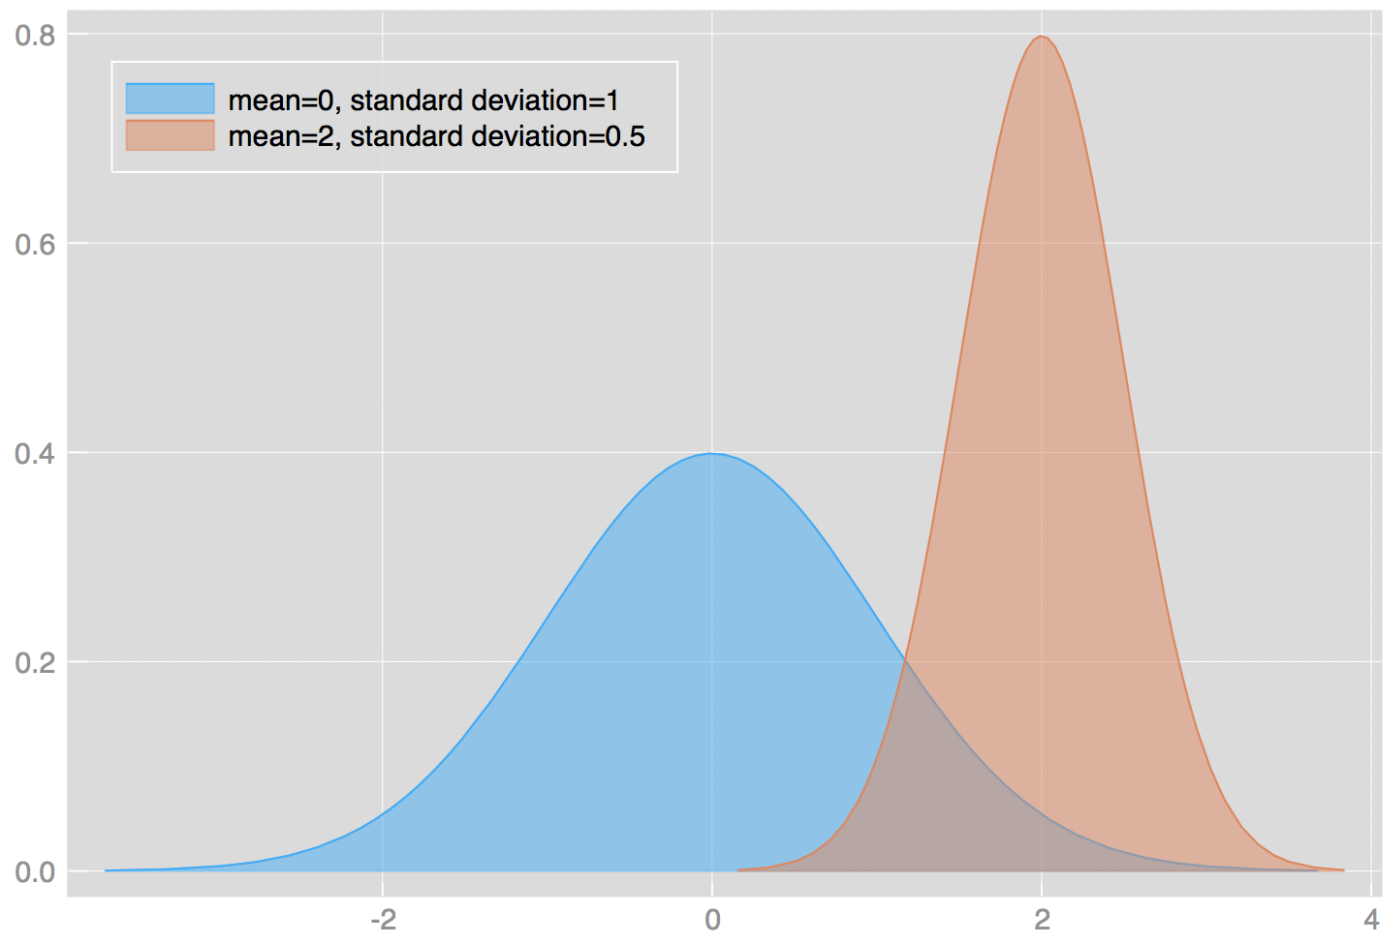
Generally, from the mean and standard deviation of the data, we can approximate the properties of the population. (Note: the modeled central tendency does not have to be the mean, and the spread does not have to be the standard deviation)

VAEs learn stochastic/probabilistic mappings between the input space and latent space

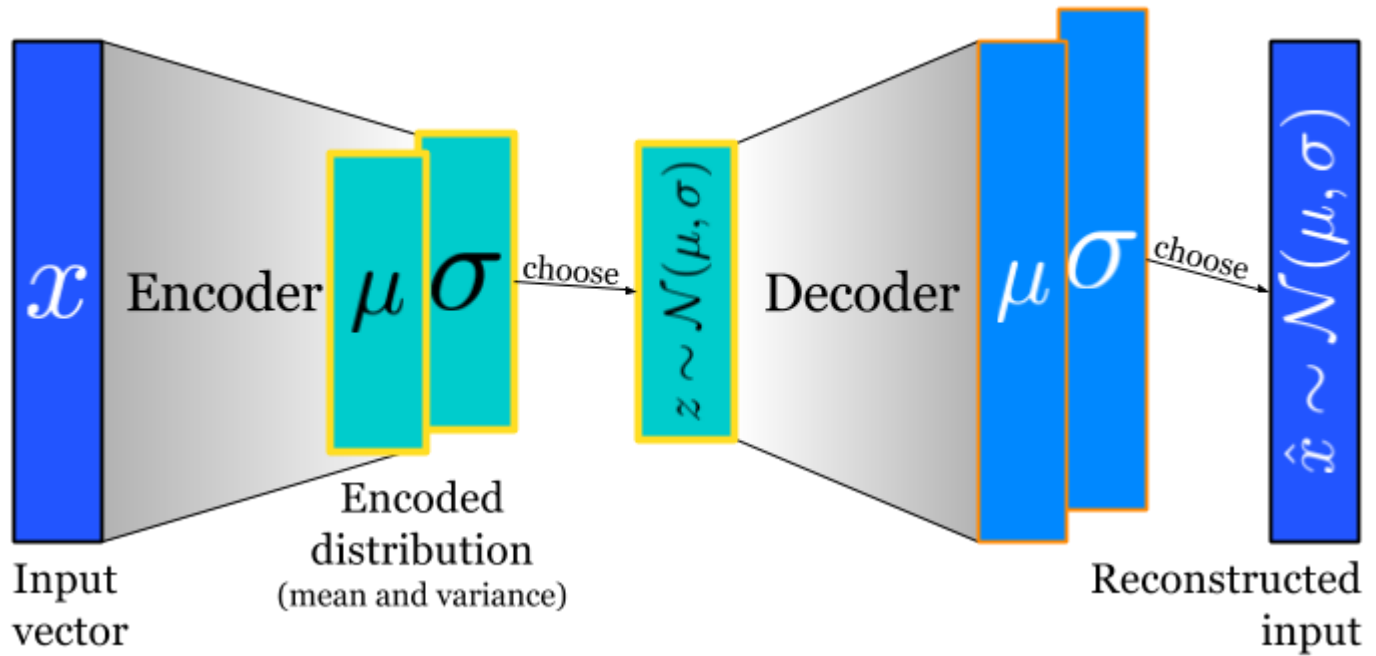
In this tutorial, the population represents all images that can be in the category of class of training data.



Latent Space: if assumed to have Gaussian distribution, it's parameters are mean, and standard deviation



# Variational Autoencoder



Variational Autoencoder loss considers two things:

1. The negative log likelihood of the output  $x_i$  multiplied by their corresponding weight (or probability)  $p_i$

$$-E_{z \sim q_\theta(z|x_i)} [\log(p_\phi(x_i|z))]$$

$E[X] = \sum_{i=1}^n x_i p_i$  expectation (i.e. expected or average value) is the weighted sum of all examples

$q_\theta(z|x_i)$  Learned latent space distribution

$p_\phi(x_i|z)$  Distribution of  $x$  given  $z$  (reconstruction distribution)

2. Kullback-Leibler Divergence (KLD) of the "actual" (i.e. the prior) distribution and the predicted distribution.

- KLD metric describing difference between two distributions
- Ideally the difference between the true distribution and the modeled distribution should as small as possible

$$KL(q_\phi(z|x_i)||p(z)) = q_\phi(z|x_i) * \log\left(\frac{q_\phi(z|x_i)}{p(z)}\right)$$

Combining parts (1) and (2) to construct our loss function:

$$loss = l_i(\theta, \phi) = -E_{z \sim q_\theta(z|x_i)} [\log(p_\phi(x_i|z))] + KL(q_\phi(z|x_i)||p(z))$$

```
1 # LOAD DATA
2 (training x, training y), (testing x, testing y) = fashion_mnist.load_data()
```

```

2 (\training_x, training_y), (\testing_x, testing_y), fashion_mnist.test_data,
3 training_x.shape, training_y.shape, testing_x.shape, testing_y.shape

((60000, 28, 28), (60000,), (10000, 28, 28), (10000,))

1 # SPLIT TRAINING SET
2 train_size = 50000
3 train_x = training_x[:train_size]
4 train_y = training_y[:train_size]
5 val_x = training_x[train_size:]
6 val_y = training_y[train_size:]
7
8 # PRE-PROCESS DATA
9 train_x = train_x / 255.
10 val_x = val_x / 255.
11 test_x = testing_x / 255.
12
13 # VECTORIZE DATA
14 train_x = train_x.reshape(-1, 784)
15 val_x = val_x.reshape(-1, 784)
16 test_x = test_x.reshape(-1, 784)
17
18 train_x.shape, val_x.shape, test_x.shape

((50000, 784), (10000, 784), (10000, 784))

1 # TODO: Input placeholder
2 input_img = Input(shape=(784,))
3
4 # Encoded input representation
5 l1_out = Dense(500, activation='relu', name='encoder_L1')(input_img)
6 z_mu = Dense(10, name='z_mu')(l1_out)
7 z_log_sigma = Dense(10, name='z_log_sigma')(l1_out)
8
9
10 # Define layer to incorporate KL divergence into the training loss
11 class KLDDLayer(Layer):
12     """
13     Layer designed to incorporate the loss associated with the
14     latent space distribution
15     """
16     def __init__(self, *args, **kwargs):
17         self.is_placeholder = True
18         super(KLDDLayer, self).__init__(*args, **kwargs)
19     def call(self, inputs):
20         mu, log_sigma = inputs
21         kl_batch = -.5 * K.sum(1 + log_sigma -
22                                K.square(mu) -
23                                K.exp(log_sigma), axis=-1)
24         kl_loss = K.mean(kl_batch)
25         self.add_loss(kl_loss, inputs=inputs)

```

```

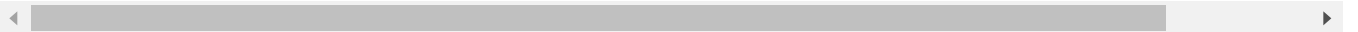
26         return inputs
27
28     # Define function for sampling in the latent space
29     def sampling(args):
30         """
31         PARAMS:
32             args = (z_mean, z_log_var)
33                 z_mean (tensor): mean of the latent space
34                 z_log_var (tensor): log of the variance of the latent space
35         RETURN:
36             z (tensor): a sample from the latent space
37         """
38         z_mu, z_log_var = args
39         batch_size = K.shape(z_mu)[0]
40         latent_dim = K.int_shape(z_mu)[1]
41         epsilon = K.random_normal(shape=(batch_size, latent_dim))
42         return z_mu + K.exp(.5 * z_log_var) * epsilon
43
44
45     # Create the KLD layer for the model
46     z_mu, z_log_sigma = KLDLayer()([z_mu, z_log_sigma])
47
48     # Sample from the latent space distribution
49     z = Lambda(sampling, output_shape=(10,))([z_mu, z_log_sigma])
50
51     # Model mapping to input representation (i.e. the encoder)
52     v_encoder = Model(input_img, z)
53
54
55     # Model mapping input to its reconstruction
56     v_decoder = Sequential([
57         Dense(500, input_dim=10, activation='relu', name='decoder_L1'),
58         Dense(784, activation='sigmoid', name='decoder_recon')
59     ])
60     recon = v_decoder(z)
61
62
63     # Model mapping an input to its reconstruction
64     vae = Model(input_img, recon)
65
66     # Display model summary
67     vae.summary()
68
69     # Function to compute the negative log likelihood
70     def nll(y_true, y_pred):
71         return K.sum(K.binary_crossentropy(y_true, y_pred), axis=-1)
72
73     # Compile model
74     vae.compile(optimizer='adam', loss=nll)

```

Model: "functional\_19"

---

Layer (type)	Output Shape	Param #	Connected to
input_2 (InputLayer)	[(None, 784)]	0	
encoder_L1 (Dense)	(None, 500)	392500	input_2[0][0]
z_mu (Dense)	(None, 10)	5010	encoder_L1[0][0]
z_log_sigma (Dense)	(None, 10)	5010	encoder_L1[0][0]
kld_layer (KLDLayer)	[(None, 10), (None, 10)]	0	z_mu[0][0] z_log_sigma[0][0]
lambda (Lambda)	(None, 10)	0	kld_layer[0][0] kld_layer[0][1]
sequential (Sequential)	(None, 784)	398284	lambda[0][0]
Total params: 800,804			
Trainable params: 800,804			
Non-trainable params: 0			



```

1  # Train the model to generate images
2  t0 = time()
3  train_history = vae.fit(train_x, train_x, epochs=20, batch_size=2048,
4                          validation_data=(val_x, val_x))
5  t1 = time()
6  etime = float(t1 - t0) / 60
7  print("Elapsed time: %.02f min" % etime)
8
9  # Plot Learning Loss
10 plot_learning_loss(train_history)

```

```

Epoch 1/20
25/25 [=====] - 5s 199ms/step - loss: 439.4274 - val_loss: 356
Epoch 2/20
25/25 [=====] - 5s 195ms/step - loss: 327.6814 - val_loss: 308
Epoch 3/20
25/25 [=====] - 5s 194ms/step - loss: 298.2952 - val_loss: 289
Epoch 4/20
25/25 [=====] - 5s 195ms/step - loss: 283.7543 - val_loss: 280
Epoch 5/20
25/25 [=====] - 5s 193ms/step - loss: 275.5267 - val_loss: 273
Epoch 6/20
25/25 [=====] - 5s 193ms/step - loss: 270.1948 - val_loss: 269
Epoch 7/20
25/25 [=====] - 5s 193ms/step - loss: 266.6684 - val_loss: 267
Epoch 8/20
25/25 [=====] - 5s 201ms/step - loss: 264.2265 - val_loss: 264
Epoch 9/20
25/25 [=====] - 5s 205ms/step - loss: 262.0324 - val_loss: 262
Epoch 10/20
25/25 [=====] - 5s 206ms/step - loss: 260.1703 - val_loss: 260
Epoch 11/20
25/25 [=====] - 5s 207ms/step - loss: 259.0315 - val_loss: 259
Epoch 12/20
25/25 [=====] - 5s 207ms/step - loss: 257.5126 - val_loss: 258
Epoch 13/20
25/25 [=====] - 5s 206ms/step - loss: 256.8858 - val_loss: 257
Epoch 14/20
25/25 [=====] - 5s 206ms/step - loss: 255.6491 - val_loss: 256
Epoch 15/20
25/25 [=====] - 5s 203ms/step - loss: 255.0017 - val_loss: 257
Epoch 16/20
25/25 [=====] - 5s 197ms/step - loss: 254.8318 - val_loss: 255
Epoch 17/20
25/25 [=====] - 5s 205ms/step - loss: 253.5309 - val_loss: 254
Epoch 18/20
25/25 [=====] - 5s 206ms/step - loss: 252.8493 - val_loss: 253
Epoch 19/20

```

```

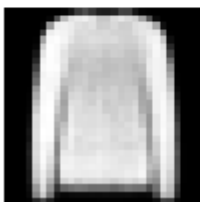
1 # RECONSTRUCTION
2 plot_compare_reconstruction(vae, test_x)

```

Reconstruction



Actual

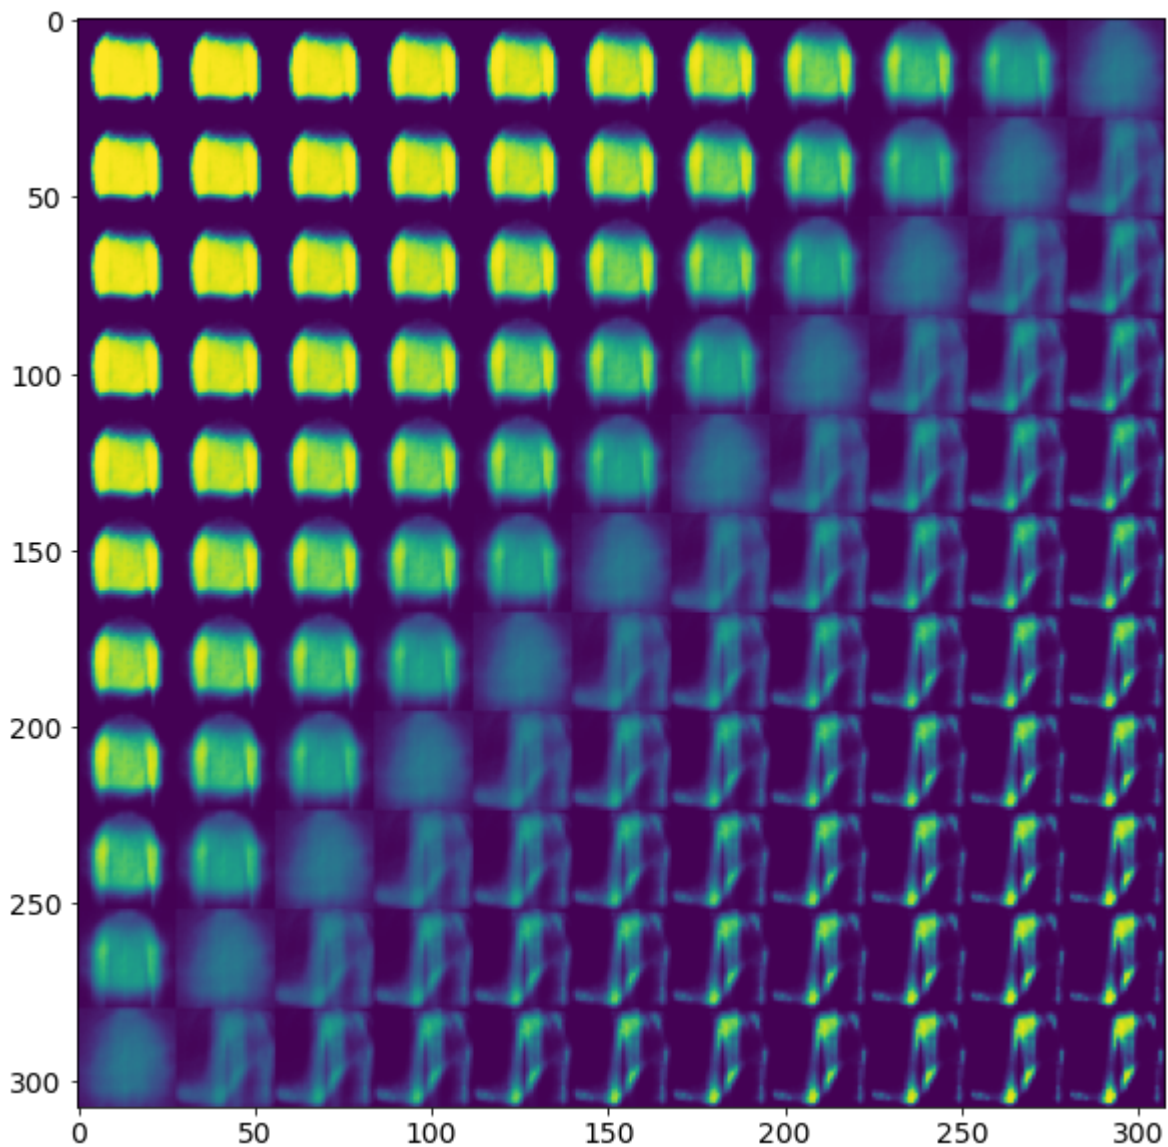


Learning Epoch

```

1  # GENERATION AND LINEAR INTERPOLATION IN THE LATENT SPACE
2  n = 11 # figure with 11x11 imgs
3  img_size = 28
4  figure = np.zeros((img_size * n, img_size * n))
5  # we will sample n points within [-11, 11] standard deviations
6  grid_x = np.linspace(-2, 2, n)
7  grid_y = np.linspace(-2, 2, n)
8
9  for i, yi in enumerate(grid_x):
10     for j, xi in enumerate(grid_y):
11         z_sample = np.repeat(np.array([[xi + yi]]), 10, axis=1)
12         x_decoded = v_decoder.predict(z_sample)
13         img = x_decoded[0].reshape(img_size, img_size)
14         figure[i * img_size: (i + 1) * img_size,
15             j * img_size: (j + 1) * img_size] = img
16
17 plt.figure(figsize=(10, 10))
18 plt.imshow(figure)
19 plt.show()

```



```

1  # DISPLAY IMAGES
2  n = 11 # figure with 11x11 imgs
3  img_size = 28
4  figure = np.zeros((img_size * n, img_size * n))
5  # we will sample n points within [-11, 11] standard deviations
6  grid_x = np.linspace(0, 5, n)
7  grid_y = np.linspace(0, 5, n)
8
9  k = 0
10 for i, yi in enumerate(grid_x):
11     for j, xi in enumerate(grid_y):
12         # Display example image
13         img = train_x[k].reshape(28,28) # i + j
14         #plt.title('Class: %d' % train_y[0])
15         figure[i * img_size: (i + 1) * img_size,
16               j * img_size: (j + 1) * img_size] = img
17         k += 1
18
19 plt.figure(figsize=(10, 10))
20 plt.imshow(figure)
21 plt.show()

```





```
1 gt1 = time()
2 g_etime = (gt1 - gt0) / 60
3 print("Global Elapsed Time: %.02f min" % g_etime)
```

Global Elapsed Time: 29.92 min



## ▼ Closing



QUESTIONS?

THANK YOU

[Post-knowledge Survey](#)

If you went through the tutorial please complete this brief [sign in form](#)

Facebook [@oucsgsa](#)



Instagram



